

Intro to Multiprocessor Programming

HomeWork1

Dror Nir 2030572989

Uria Mor 200672954

Q1

Assume we have 3 threads A , B , C and they all try to enter level 1 together at the following order:

$$\begin{aligned} &\text{write}_A(\text{level}[A] = 1) \rightarrow \text{write}_A(\text{victim}[1] = A) \rightarrow \text{write}_B(\text{level}[B] = 1) \rightarrow \text{write}_B(\text{victim}[1] = B) \rightarrow \\ &\quad \rightarrow \text{write}_C(\text{level}[C] = 1) \rightarrow \text{write}_C(\text{victim}[1] = C) \end{aligned}$$

so clearly, $D_A \rightarrow D_B$, now we can examine the following possible chain of events:

$$\begin{aligned} &\rightarrow \text{read}_B(\text{victim}[1] = C) \rightarrow \text{write}_B(\text{level}[B] = 2) \rightarrow \text{write}_B(\text{victim}[2] = B) \rightarrow \\ &\quad \rightarrow \text{read}_B(\text{level}[A] = 1) \rightarrow \text{read}_B(\text{level}[C] = 1) \end{aligned}$$

since all threads but B are at level ≤ 2 , it's possible for B to enter it's critical section:

$$\rightarrow CS_B \rightarrow \text{write}_B(\text{level}[B] = 0)$$

It's only natural to have B trying to enter level 1 again:

$$\rightarrow \text{write}_B(\text{level}[B] = 1) \rightarrow \text{write}_B(\text{victim}[1] = B)$$

Notice that nothing keeps C from reading $\text{victim}[1] = B$ before A does, so it's possible to continue with

$$\rightarrow \text{read}_C(\text{victim}[1] = B)$$

and we can let C reach (and finish) it's CS before A reads $\text{victim}[1] \neq A$:

$$\rightarrow CS_C \rightarrow \text{write}_C(\text{level}[C] = 0)$$

In the same manner as above we can send B and C interchangeably to finish an unbounded number of “re-wind $\rightarrow \text{level-up} \rightarrow CS$ ” execution cycles before having A reaching the event $\text{read}_A(\text{victim}[1] \neq A)$. So although $D_A \rightarrow D_B$, we get that for any $k \geq 0$ it's possible for $CS_B^{1+k} \rightarrow CS_A$ to occur.

Q2

Lemma. *Two threads performing only their first event will cause the shared memory state to be the same, regardless of order*

Proof. *If The first operation in $\text{lock}()$ is a read: reading obviously can't change the shared memory state.*

If The first operation in $\text{lock}()$ is a write: because it is the first operation, it means that the data that is written is constant, and it always writes to the same register. This is because the thread didn't read anything, including any information that can cause it to change the data written or location to write to. Therefore, order of writing has no effect.

We will let $T1$ perform the first action, then $T2$ will perform the first action. Then, we will let one of the threads run until it unlocks (It is deadlock-free so this must happen). We claim that regardless of the choice of thread to run, it will behave the same as the other thread will behave. This is because the shared memory state is identical regardless of choice of the first thread (Lemma). Therefore, it doesn't matter which thread started first, a contradiction to strong FCFS.

Q3

1. Normalizing the (total) sequential execution time of the program to be 1, while the execution time is comprised out of M which is 40% of the total time - $\frac{2}{5}$ then the unparallelizable portion of the program is $1 - p = \frac{2}{5}$ thus $p = \frac{3}{5}$. Using Amdahl's formula we get

$$\begin{aligned} S &= \frac{1}{\frac{2}{5} + \frac{3/5}{n}} \\ &= \frac{1}{\left(\frac{1}{5}\right)\left(2 + \frac{3}{n}\right)} \\ &= \frac{5n}{2n+3} \end{aligned}$$

which gives us an upper bound of 2.5 folds of overall speedup.

2. Assume the sequential execution time of the program is T . Since the rest of the program is completely parallelizable, we can express the total running time on n processors as $\frac{3}{10}T + \frac{7T}{n}$.
denote M 's speedup by S' , then $\frac{3T}{S'} + \frac{7T}{n} = \frac{3T + \frac{7T}{n}}{2}$.

$$\begin{aligned} \frac{T}{10} \left(\frac{3}{S'} + \frac{7}{n} \right) &= \frac{T}{10} \frac{3 + \frac{7}{n}}{2} \\ &\Downarrow \\ \frac{3}{S'} + \frac{7}{n} &= \frac{3 + \frac{7}{n}}{2} \\ &\Downarrow \\ \frac{3}{S'} &= \frac{3 - \frac{7}{n}}{2} \\ &\Downarrow \\ S' &= \frac{3 \cdot 2}{3 - \frac{7}{n}} = \frac{6n}{3n - 7} \end{aligned}$$

3. Denote M 's part in the overall execution time by $1 - p$, given that $\frac{1}{\frac{1-p}{3} + \frac{p}{n}} = \frac{2}{1-p+\frac{p}{n}}$ (twice the speedup compared to running on n processors without M being sped up), solution is $1 - p = \frac{3}{n+3}$.

Q4

1. Without loss of generality, B finished the `lock()` code after A did (finishing is running Line 11 when `turn == me`). Let's mark the last time thread A executed Line 8 as t_1 , and similarly t_2 for B. Because B finished after A, $t_1 < t_2$. Because t_1 is the last time A executed Line 8, A must finish Line 9 before B runs Line 10. So in the (immediately) next time A runs Line 11 "turn" is not equal to "me", and A runs Line 8 one more time - a contradiction for t_1 being the last time A ran Line 8.
2. Counter example: A sets "turn", A sets "busy" to TRUE, B sets "turn", A checks if it's his turn (it's not) so it loops in the inner "while" until someone else calls `unlock()`. However, B is now also stuck in the loop because "busy" is true. And actually any other thread that wants to lock will get stuck in the loop, and no-one can finish the `lock()`, so no-one can `unlock()`.
3. Starvation-free implies deadlock-free.

Q5

We offer two solutions, the first is a reduction from the *MRMW* case, and the second is basically common sense. Since we'd very much like to understand why the later solution doesn't work, we submit both of them.


1. We will prove the same theorem, but for *MRMW* registers. Then it will obviously hold for *MRSW* registers. In class, we have shown that for $k = 2$, the statement holds.

Assuming there is a covering state where $k - 1$ threads not in the CS cover $k - 1$ distinct registers: WLOG there exists some thread t that did not run yet, t is not in the set of $k - 1$ threads that are in the covering state and t is also not in the CS. We let t run until it's covering a new register or is in the CS, whichever comes first (note we can always go into the CS because of deadlock-freedom). If t covered a new register, we're done. Otherwise, t might have written to any of the $k - 1$ registers already covered, but we can "release" all the $k - 1$ threads and let them run and immediately overwrite what t had written. Consequently, The state of memory that is visible to the $k - 1$ threads is identical to the state before t ran, so at least one thread out of the $k - 1$ threads will enter the CS (deadlock-free), a contradiction to mutual exclusion.


Therefore, t must have entered a covering state on a new register before reaching the CS.

2. If we have k *MRSW* registers and N threads, if $k < N$, then by the pigeonhole principle, we have at least two threads that will write to the same register before entering CS. By the definition of *MRSW*, this is not possible, thus, at least N *MRSW* registers in a dead-lock free mutex algorithm.

Q6


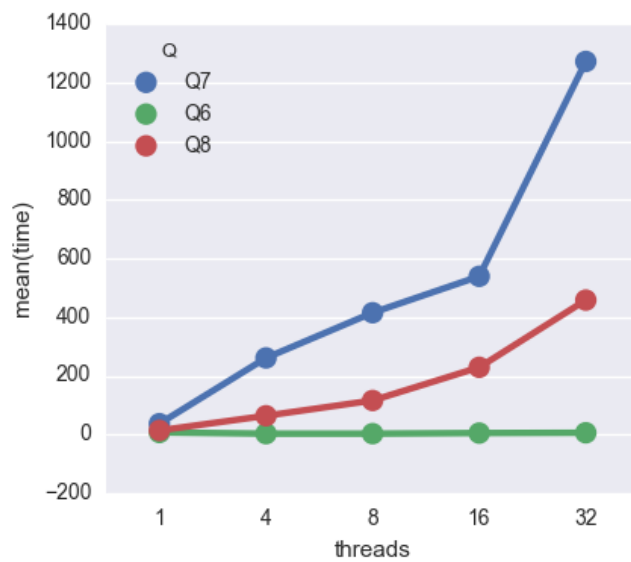

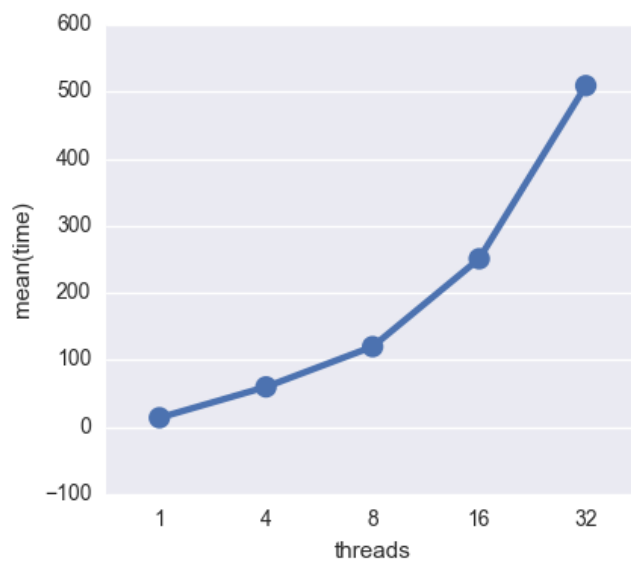


Q7



First, we see that the counter was locked appropriately: its value is the correct value as expected, in comparison to the previous question. Now we look at the graph and see that it appears to have non-linear growth. The slowdown is because threads must wait for the lock to be unlocked. This result was expected because of Amdahl's Law: Most of our program is not parallelized

Q8



On the bottom right plot, question 7's line is masked by question 8's line since the counter values are the same.